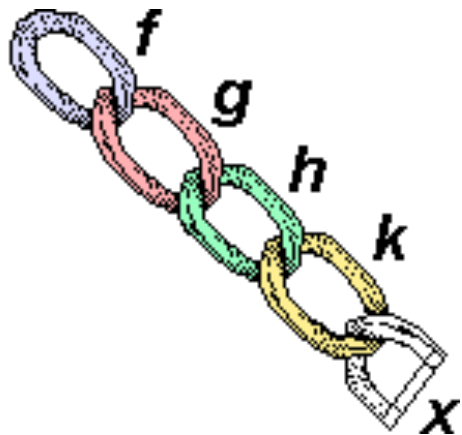




动手学深度学习 v2

李沐 · Amazon 资深首席科学家

链式法则和自动求导



Generalize to Vectors



- Chain rule for scalars:

$$y = f(u), u = g(x) \quad \frac{\partial y}{\partial x} = \frac{\partial y}{\partial u} \frac{\partial u}{\partial x}$$

- Generalize to vectors straightforwardly

$$\begin{array}{ccc} \frac{\partial y}{\partial \mathbf{x}} = \frac{\partial y}{\partial u} \frac{\partial u}{\partial \mathbf{x}} & \frac{\partial y}{\partial \mathbf{x}} = \frac{\partial y}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}} & \frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial \mathbf{y}}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}} \\ (1,n) \quad (1,) \quad (1,n) & (1,n) \quad (1,k) \quad (k,n) & (m,n) \quad (m,k) \quad (k,n) \end{array}$$

Example 1

$$\frac{\partial y}{\partial \mathbf{x}} = \frac{\partial y}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$$

Assume $\mathbf{x}, \mathbf{w} \in \mathbb{R}^n$, $y \in \mathbb{R}$

$$z = (\langle \mathbf{x}, \mathbf{w} \rangle - y)^2$$

Compute $\frac{\partial z}{\partial \mathbf{w}}$

$$a = \langle \mathbf{x}, \mathbf{w} \rangle$$

$$b = a - y$$

$$z = b^2$$

Decompose

$$\begin{aligned} \frac{\partial z}{\partial \mathbf{w}} &= \frac{\partial z}{\partial b} \frac{\partial b}{\partial a} \frac{\partial a}{\partial \mathbf{w}} \\ &= \frac{\partial b^2}{\partial b} \frac{\partial a - y}{\partial a} \frac{\partial \langle \mathbf{x}, \mathbf{w} \rangle}{\partial \mathbf{w}} \\ &= 2b \cdot 1 \cdot \mathbf{x}^T \\ &= 2 (\langle \mathbf{x}, \mathbf{w} \rangle - y) \mathbf{x}^T \end{aligned}$$

Example 2

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial \mathbf{y}}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$$

Assume $\mathbf{X} \in \mathbb{R}^{m \times n}$, $\mathbf{w} \in \mathbb{R}^n$, $\mathbf{y} \in \mathbb{R}^m$

$$z = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

Compute $\frac{\partial z}{\partial \mathbf{w}}$

$$\begin{aligned} \frac{\partial z}{\partial \mathbf{w}} &= \frac{\partial z}{\partial \mathbf{b}} \frac{\partial \mathbf{b}}{\partial \mathbf{a}} \frac{\partial \mathbf{a}}{\partial \mathbf{w}} \\ &= \frac{\partial \|\mathbf{b}\|^2}{\partial \mathbf{b}} \frac{\partial \mathbf{a} - \mathbf{y}}{\partial \mathbf{a}} \frac{\partial \mathbf{X}\mathbf{w}}{\partial \mathbf{w}} \\ &= 2\mathbf{b}^T \times \mathbf{I} \times \mathbf{X} \\ &= 2(\mathbf{X}\mathbf{w} - \mathbf{y})^T \mathbf{X} \end{aligned}$$

Decompose

$$\begin{aligned} \mathbf{a} &= \mathbf{X}\mathbf{w} \\ \mathbf{b} &= \mathbf{a} - \mathbf{y} \\ z &= \|\mathbf{b}\|^2 \end{aligned}$$

Auto Differentiation (AD)



- AD evaluates gradients of a function specified by a program at given values
- AD differs to

- Symbolic differentiation

```
In[1]:= D[4 x^3 + x^2 + 3, x]
```

```
Out[1]= 2 x + 12 x^2
```

- Numerical differentiation

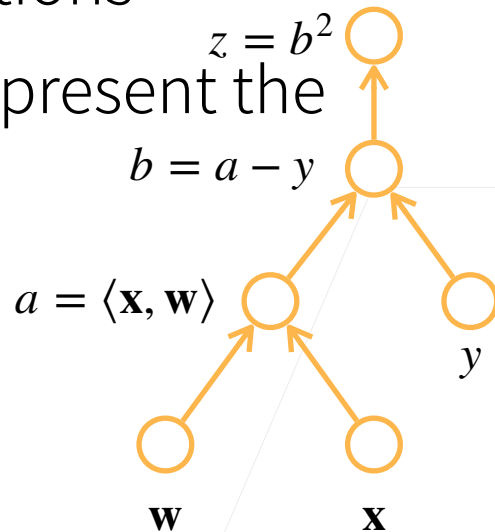
$$\frac{\partial f(x)}{\partial x} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

Computation Graph



Assume $z = (\langle \mathbf{x}, \mathbf{w} \rangle - y)^2$

- Decompose into primitive operations
- Build a directed acyclic graph to present the computation





Computation Graph

- Decompose into primitive operations
- Build a directed acyclic graph to present the computation
- Build explicitly
 - Tensorflow/Theano/MXNet

```
from mxnet import sym
a = sym.var()
b = sym.var()
c = 2 * a + b
# bind data into a and b later
```

Computation Graph



- Decompose into primitive operations
- Build a directed acyclic graph to represent computation
 - from mxnet import autograd, nd
 - with autograd.record():
 - a = nd.ones((2,1))
 - b = nd.ones((2,1))
 - c = 2 * a + b
- Build explicitly
 - Tensorflow/Theano/MXNet
- Build implicitly though tracing
 - PyTorch/MXNet



Two Modes

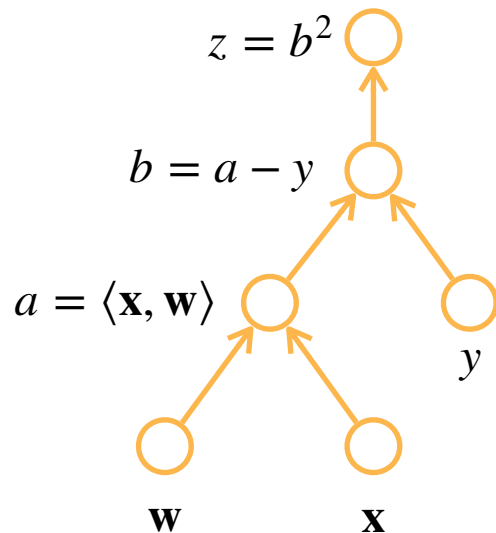
- By chain rule
$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial u_n} \frac{\partial u_n}{\partial u_{n-1}} \dots \frac{\partial u_2}{\partial u_1} \frac{\partial u_1}{\partial x}$$
- Forward accumulation
$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial u_n} \left(\frac{\partial u_n}{\partial u_{n-1}} \left(\dots \left(\frac{\partial u_2}{\partial u_1} \frac{\partial u_1}{\partial x} \right) \right) \right)$$
- Reverse accumulation (a.k.a Backpropagation)
$$\frac{\partial y}{\partial x} = \left(\left(\left(\frac{\partial y}{\partial u_n} \frac{\partial u_n}{\partial u_{n-1}} \right) \dots \right) \frac{\partial u_2}{\partial u_1} \right) \frac{\partial u_1}{\partial x}$$

Reverse Accumulation

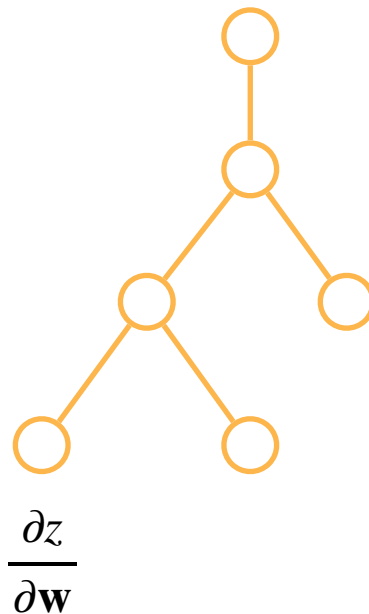


Assume $z = (\langle \mathbf{x}, \mathbf{w} \rangle - y)^2$

Forward



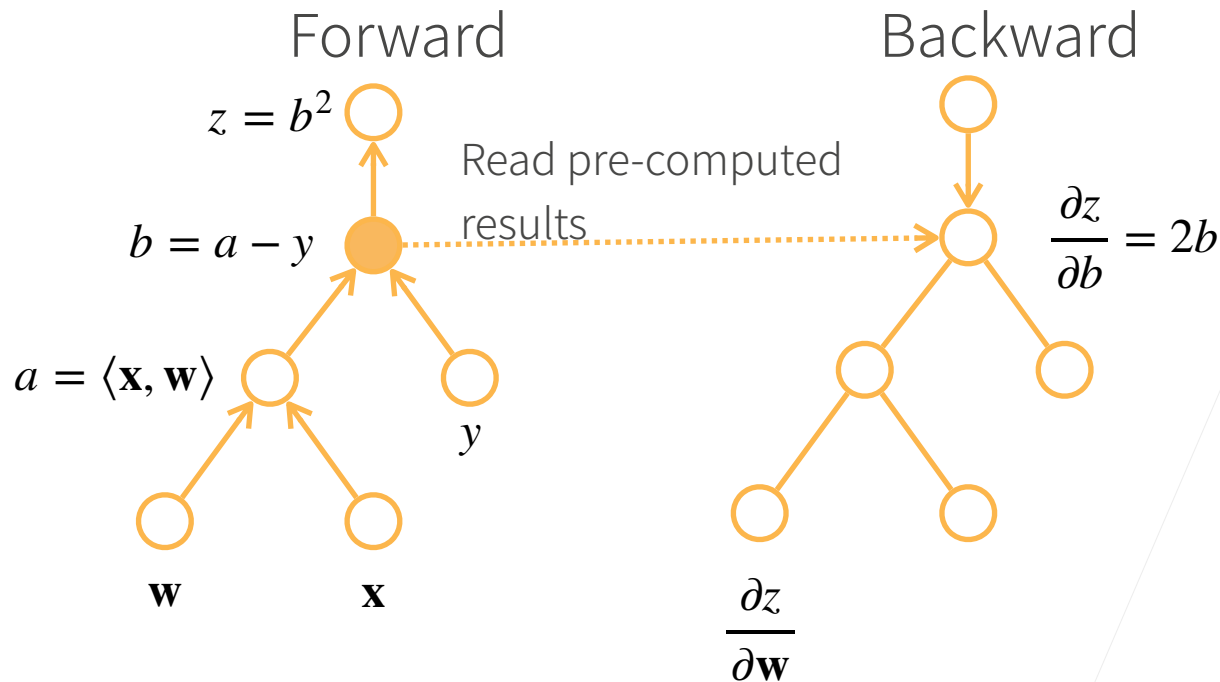
Backward



Reverse Accumulation



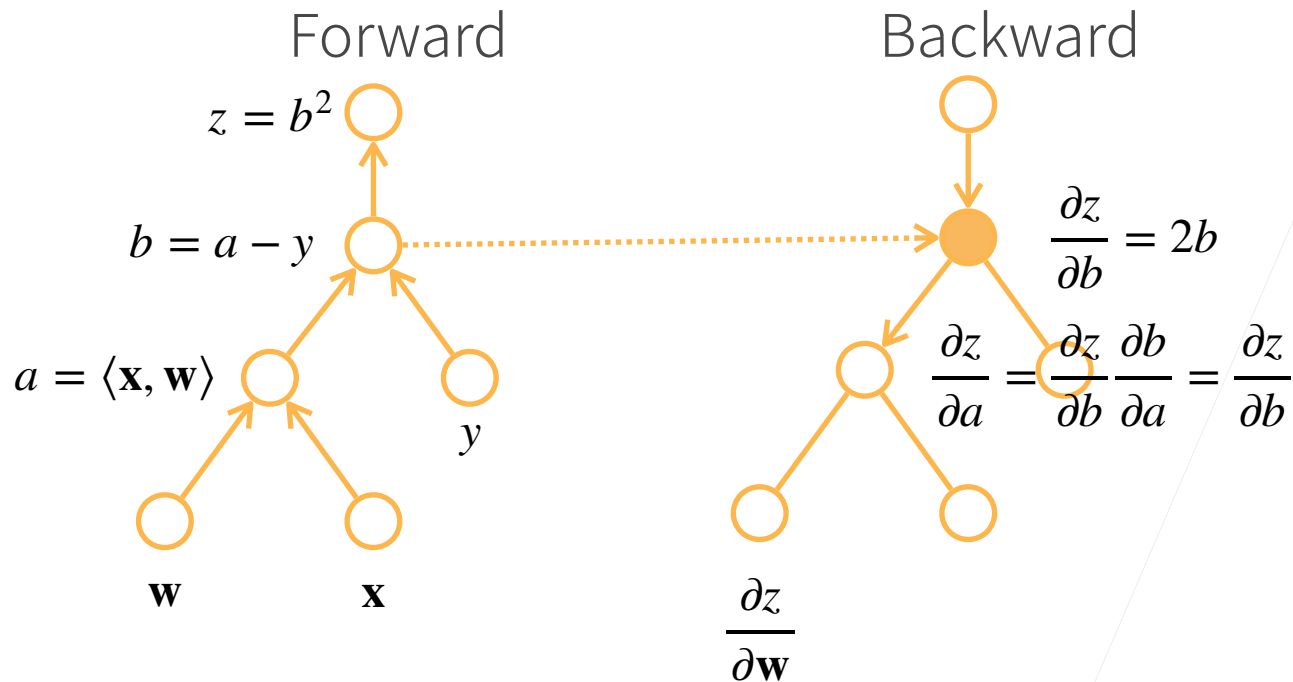
Assume $z = (\langle \mathbf{x}, \mathbf{w} \rangle - y)^2$



Reverse Accumulation



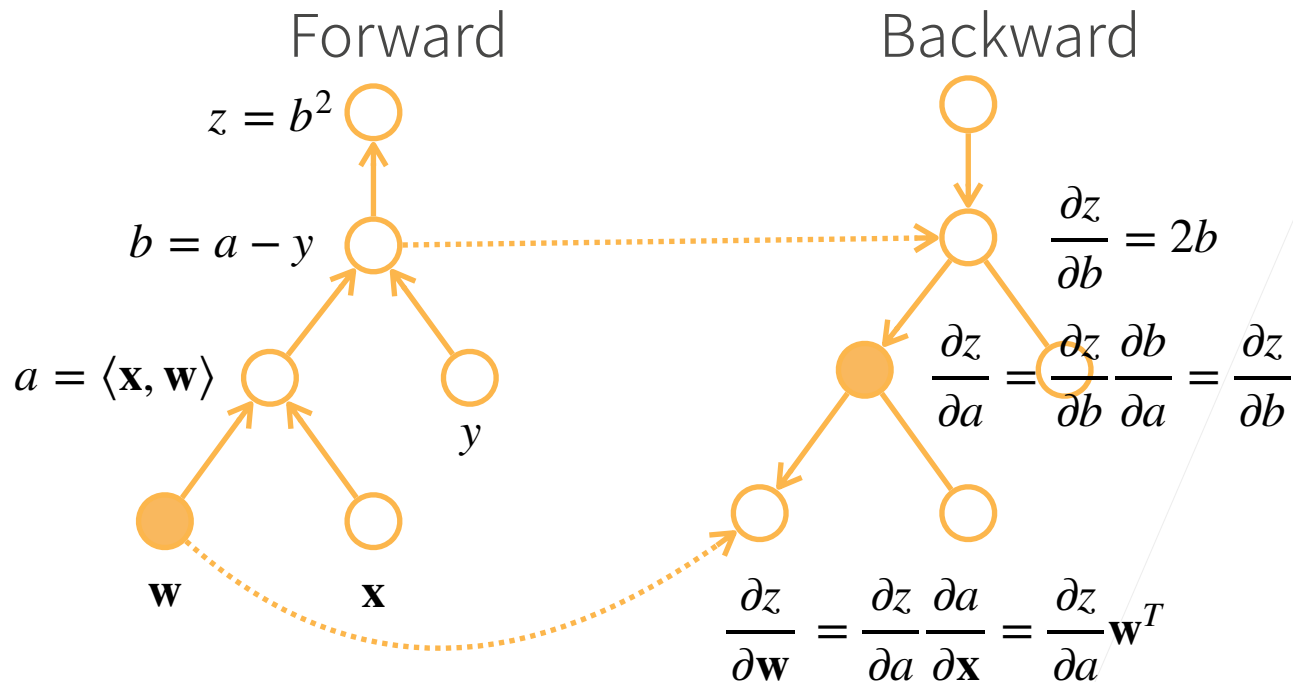
Assume $z = (\langle \mathbf{x}, \mathbf{w} \rangle - y)^2$



Reverse Accumulation



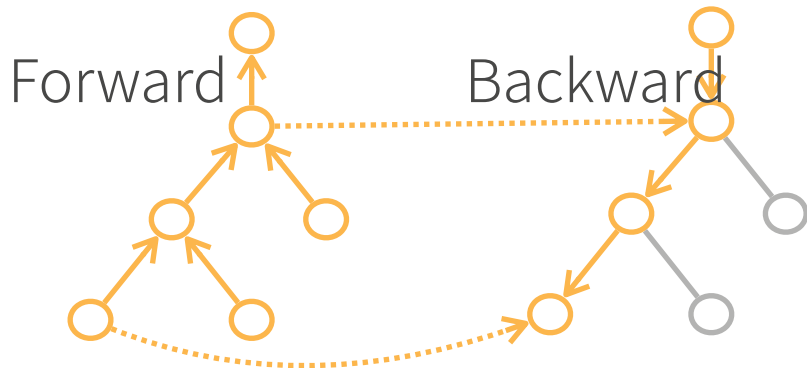
Assume $z = (\langle \mathbf{x}, \mathbf{w} \rangle - y)^2$



Reverse Accumulation Summary



- Build a computation graph
- Forward: Evaluate the graph, store intermediate results
- Backward: Evaluate the graph in a reversed order
 - Eliminate paths not needed



Complexities



- Computational complexity: $O(n)$, n is #operations, to compute all derivatives
 - Often similar to the forward cost
- Memory complexity: $O(n)$, needs to record all intermediate results in the forward pass
- Compare to forward accumulation:
 - $O(n)$ time complexity to compute one gradient, $O(n \cdot k)$ to compute gradients for k variables
 - $O(1)$ memory complexity



[Advanced] Rematerialization

- Memory is bottleneck for backward accumulation
 - Linear to #layers and batch size
 - Limited GPU memory (32GB max)
- Trade computation for memory
 - Save a part of intermediate results
 - Recompute the rest when needed

Rematerialization

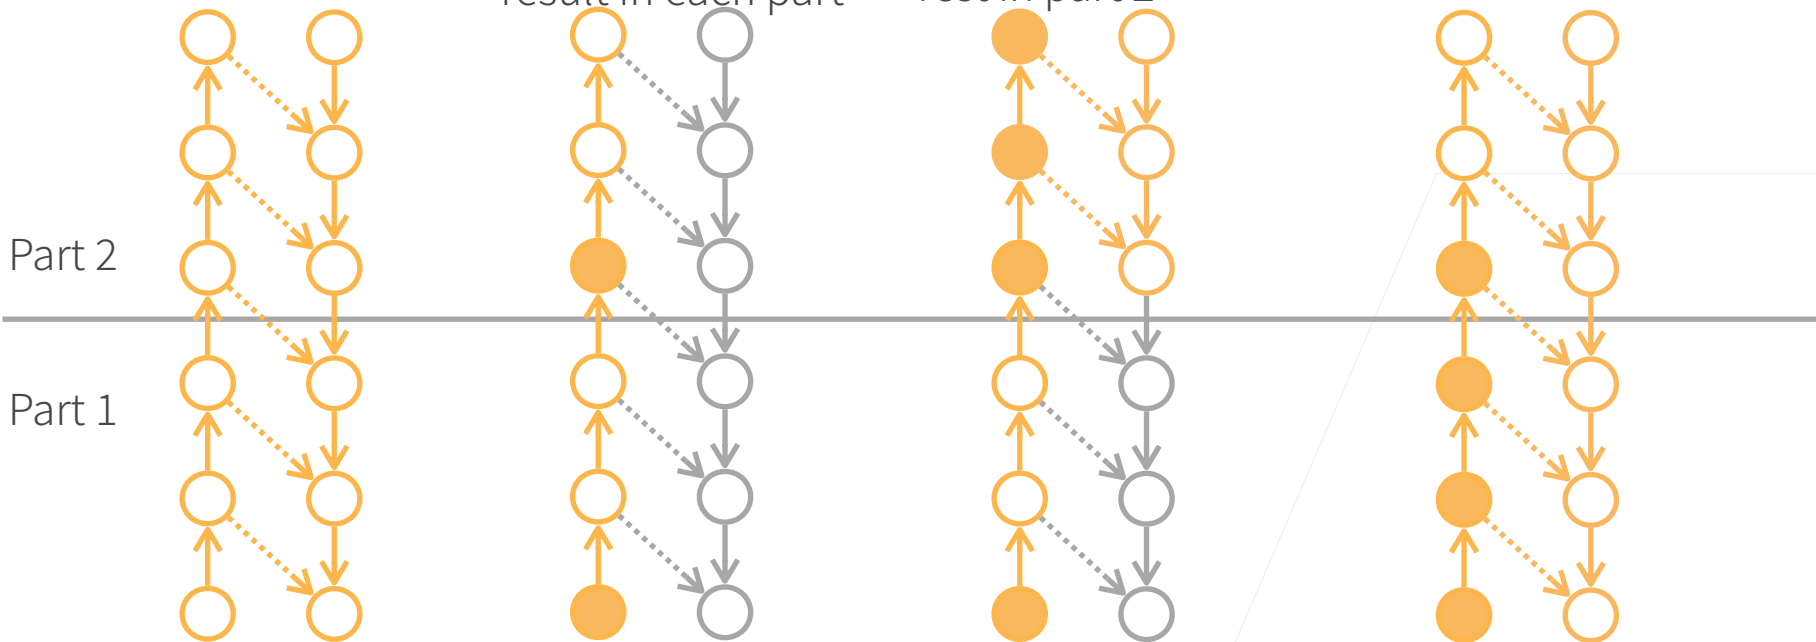


Forward Backward

Only store the head
result in each part

Recompute the
rest in part 2

Recompute the
rest in part 1



Complexities



- An additional forward pass
- Assume m parts, then $O(m)$ for head results, $O(n/m)$ to store one part's results

- Choose $m = \sqrt{n}$ then the memory complexity is

$$O(\sqrt{n})$$

- Applying to deep neural networks
 - Only throw away simple layers, e.g. activation, often <30% additional overhead
 - Train 10x larger networks, or 10x large batch size